

Project milestone - group #56

Project Title

Spoiler Alert: Who Makes it to the Finale?

Team Members Info

- Stav Abukarat, 322897687, stava, stav.abukarat@mail.huji.ac.il.
- Shahar Spencer, 208067017, shahar.spencer, shahar.spencer@mail.huji.ac.il.
- Bar Dalal, 212622260, bar.dalal, bar.dalal@mail.huji.ac.il.
- Ido Oron, 208733634, ido_oron, ido.oron@mail.huji.ac.il.

Problem description

Fictional character deaths are a central storytelling device, yet they are rarely random. Writers make deliberate choices which characters survive and which don't, based on factors such as their gender, prominence, role, and the franchise's genre and tone. This raises an interesting question: can we predict whether a fictional character will die, using information available from their wiki page and the language they use throughout the story?

Our goal is to build a binary classifier that predicts `is_dead` (1 = dead, 0 = alive) using structured character metadata, including gender, species, affiliation, wiki page length, infobox field count, franchise, content rating, etc. Beyond the classifier itself, we're interested in which features are most predictive whether death is driven more by who a character is, or by which franchise they belong to. In addition, we'll analyze the franchise scripts using text-mining techniques in order to identify the language patterns associated with a character's eventual fate.

Data


We scraped character data from 38 Fandom Wikis, including Game of Thrones, MCU, Star Wars, Harry Potter, Breaking Bad, and 33 other franchises, using the MediaWiki API. For each character page, we extracted the structured infobox, which contains biographical fields like gender, species, affiliation, date of death, cause of death. These fields were standardized across franchises and transformed into a unified tabular schema. The target label, `is_dead`, was inferred from the character's status field and the presence of a date-of-death entry.

We ended up with 23MB of data - 67,327 characters from 38 franchises, each labeled dead or alive (about 27% of them are dead). The base came from scraping Fandom wikis, and we joined in extra fields from a few public APIs and Kaggle datasets, so each character has roughly 120 columns. It's important to note that Star Wars is by far the biggest wiki, so it alone is 42,271 rows — about 63% of the data. The next largest are Harry Potter (5,121), the MCU (4,957), and Indiana Jones (2,134). Due to this imbalance, we plan to use franchise-stratified sampling during modeling, so it doesn't swamp the classifier.

In addition to the character dataset, we collected script data for multiple movies and TV series, primarily from Kaggle. For instance, the dataset of "Game of Thrones" contains eight folders, one for each season, in which there are separate text files for individual episodes. This part of the dataset weighs approximately 5MB.

Visualizations

- We think that men would die more than women. We can see that men do indeed die more, both globally and per franchise. Raw: men 30.0% vs women 21.6%. Additionally, men die more in 26 of 27 franchises (only Twilight reverses, by 0.3pp), with a mean within-franchise gap of +12.3 points.



- The figure below presents eight bar charts, one for each season of “Game of Thrones”, showing the number of words in each episode’s script. Most of the visualization makes sense - it’s reasonable that a one-hour episode contains several thousand words, and the variation in script length across episodes appears plausible. The number of episodes per season is also correct and was manually verified. However, it seems that the data on episodes 3 and 5 in season 8 is incorrect. A manual inspection of the corresponding files revealed that the scripts for these episodes were not available and that the files contain placeholder text indicating that the scripts had not been released...

